

Reinforcement Learning for LTL_f/ LDL_f : Theory and Implementation



SAPIENZA
UNIVERSITÀ DI ROMA

Marco Favorito

M.Sc. in
Engineering in Computer Science
at Sapienza, University of Rome


A.Y. 2017/2018

Introduction

- Classic Reinforcement Learning:
 - An *agent* interacts with an *environment* by taking *actions* so to maximize *rewards*;
 - No knowledge about the transition model, but assume Markov property (history does not matter);
 - Based on Markov Decision Process (MDP)
- RL for Non-Markovian Decision Process (NMRDP):
 - Rewards depend from history, not just the last transition;
 - Specify proper behaviours by using temporal logic formulas;
 - Reduce the problem to MDP (with extended state space)
- In (Brafman et al. 2018) specify reward using:
 - Linear-time Temporal Logic on Finite Traces LTL_f
 - Linear-time Dynamic Logic on Finite Traces LDL_f

Example of NMRDP: BREAKOUT

- Non-Markovian reward: remove columns from left to right
- The solution space is restricted
- Policy must depend from a sequence of states: $\rho : S^* \rightarrow A$



Goals of the thesis

- Theoretical foundations for RL over NMRDP with LTL_f/LDL_f rewards by leveraging (Brafman et al. 2018)
- Formalization and solution of a new problem: RL for LTL_f/LDL_f goals
 - two-fold representation of the world
 - Low-level, used by the learning agent
 - High-level, used to specify the goal formulas
- Deal with sparse rewards and design a way to improve exploration
- Implementation of the topics described above

LTL_f and LDL_f (De Giacomo and Vardi, 2013)

- Linear Temporal Logic on finite traces: LTL_f
 - Exactly the same syntax of LTL
 - Interpreted over finite traces
 - Next: $\bigcirc happy$
 - Eventually: $\Diamond rich$
 - Until: $reply \mathcal{U} acknowledge$
 - Always: $\Box safe$
- Linear Dynamic Logic on finite traces: LTL_f
 - Merging of LTL_f and Regular Expressions
 - $[true^*](safe)$
 - $\langle A^*; B \rangle End$
- Reasoning in LTL_f/ LDL_f :
 - transform formulas φ into NFAs or DFAs \mathcal{A}_φ
 - For every trace π and LTL_f/ LDL_f formula φ :

$$\pi \models \varphi \iff \pi \in \mathcal{L}(\mathcal{A}_\varphi)$$


From LTL_f/LDL_f formulas to automata: examples

E.g. For $\pi = \langle \{\}, \{A\}, \{A, B\} \rangle$,


$$\pi \not\models \Box A$$

$$\pi \models \Diamond A$$


$$\pi \not\models \langle (A; B)^* \rangle End$$



$\Box A$
Always A



$\Diamond A$
Eventually A



$\langle (A; B)^* \rangle End$
Alternating A and B

FLLOAT: From LTL_f/LDL_f to Automata

- Python package supporting:
 - LTL_f/LDL_f formulas: parsing, syntax and semantics
 - Translation to automata (NFA, DFA, on-the-fly DFA)
- $\Box A$:

```
from flloat.parser.ltlf import LTLfParser
parser = LTLfParser()
always_A = parser("G(A)")
dfa = always_A.to_automaton(determinize=True)
dfa.to_dot("always_A.svg")
```

- $\langle(A; B)^*\rangle End$:

```
from flloat.parser.ldlf import LDLfParser
parser = LDLfParser()
alternating_AB = parser("<(A;B)*>end")
dfa = alternating_AB.to_automaton(determinize=True)
dfa.to_dot("alternating_AB.svg")
```

RL for NMRDP with LTL_f / LDL_f rewards

- Given an NMRDP \mathcal{N} with:
 - rewards specified by LTL_f / LDL_f formulas
 - over a set of propositional symbols \mathcal{P}
 - agent's state space: $S \subseteq 2^{\mathcal{P}}$
- We can transform it into an equivalent MDP \mathcal{M} (Brafman et al. 2018)
 - the state space of \mathcal{M} is extended:

$$S' = S \times Q_1 \times \cdots \times Q_m$$

where Q_i is the set of states of \mathcal{A}_{φ_i}


- reward is given iff q_i is an accepting state of \mathcal{A}_{φ_i}
- An optimal policy for \mathcal{M} is optimal for \mathcal{N}
- We reduced RL for \mathcal{M} to RL for \mathcal{N}

RL for LTL_f/LDL_f goals: a new problem

Two-fold representation of the world \mathcal{W} :

- An agent learning an MDP with **low-level features** S , trying to optimize reward R
- LTL_f/LDL_f goals $\{(\varphi_i, r_i)_{i=1}^m\}$ over a set of **high-level features** \mathcal{F} , yielding a set of fluents configurations $\mathcal{L} = 2^{\mathcal{F}}$

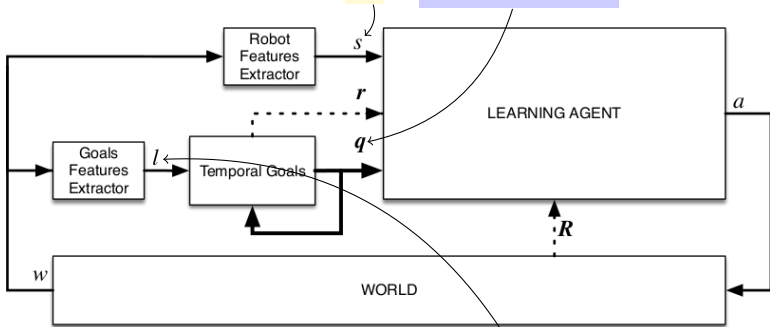
Solution: a non-Markovian policy $\rho : S^* \rightarrow A$ that is optimal wrt rewards r_i and R .



Our approach:

- Transform each φ_i into DFA \mathcal{A}_{φ_i}
- Do RL over an MDP \mathcal{M}' with a transformed state space:

$$S' = S \times Q_1 \times \dots \times Q_m$$



Notice: **the agent ignores the fluents \mathcal{L} !**

The actual RL relies on standard RL algorithms (e.g. Sarsa(λ))

Automata-based Reward Shaping

- *Reward sparsity* is the main issue: **Hard to learn complex behaviors without heuristics**

Idea: Give additional reward if, after a transition on \mathcal{A}_φ , we are closer to an accepting state (anticipate rewards)

Based on *Potential-Based Reward Shaping* (Ng et al., 1999):

$$F(s, a, s') = \gamma\Phi(s') - \Phi(s)$$

Two modes:

- *Off-line:* $\Phi(q)$ inversely proportional to the distance from any accepting state of \mathcal{A}_φ
- *On-line:* \mathcal{A}_φ and Φ are built from scratch and updated while learning and discovering new states/transitions. Using *Dynamic PBRS* (Devlin, 2012)

RLTG (Reinforcement Learning for Temporal Goals)

Reinforcement Learning Python framework that implements our approach

- Depends on flloot for the construction of \mathcal{A}_{φ_i}
- Works also for classic RL
- Highly customizable, uses OpenAI Gym interface

Example (in pseudocode):


```
env = GymEnvironment()           # a Gym environment
agent = TGAgent(                 # the "temporal goal" agent
    FeatureExtractor(...),      # generates the agent's space
    Brain(...),                # abstraction of RL algorithms (e.g. Sarsa)
    [TemporalGoal( $\varphi, r$ ), ...] # list of temp. goal managers
)
trainer = TGTrainer(env, agent, stop_conditions=..., ...)
trainer.main(...)                # starts the learning process
```

Experiments: BREAKOUT

BREAKOUT: remove columns/rows in a given order

- **low-level features:** paddle position, ball speed/position;
- **high-level features:** bricks status (broken/not broken)
- **LT_L_f/LDL_f goal** (l_i means: the i_{th} line has been removed.):

$$\langle (\neg l_0 \wedge \neg l_1 \wedge \neg l_2)^*; (l_0 \wedge \neg l_1 \wedge \neg l_2); (l_0 \wedge \neg l_1 \wedge \neg l_2)^*; \dots; (l_0 \wedge l_1 \wedge l_2) \rangle tt$$



Experiments: SAPIENTINO

SAPIENTINO: make a *bip* in each color, in a given order

- **low-level features:** robot position (x, y) ;
- **high-level features:** color of current cell, bip-last-action
- **LT_L_f/LDL_f goal:**

$$\langle (\neg \text{bip})^*; \text{red} \wedge \text{bip}; (\neg \text{bip})^*; \text{green} \wedge \text{bip}; \dots \rangle tt$$




Experiments: MINECRAFT

MINECRAFT: complete tasks by getting resources/using tools

- **low-level features:** robot position (x, y) ;
- **high-level features:** *get_wood*, *use_workbench* etc.
- LTL_f/LDL_f **goal:** three composite tasks like (approximately):

$\langle true^* \rangle \langle get_iron; get_iron \wedge get_wood; \dots \wedge use_factory \rangle tt$



Discussion

Why a two-fold representation for temporal goals?

- *Separation of concerns*: It allows to a better design of the RL system by improving modularity and flexibility;
- *Reduced state space*: It makes easier the learning by exploring a smaller state space and focusing only on relevant low-level features;
- *Simpler agent*: Move part of the complexity outside the agent.

Other observations:

- *Expressive Power*: LDL_f expressive as Monadic Second-Order logic (e.g. we can specify procedural constraints);
- No need of new algorithms, one can rely on off-the-shelf RL algorithms (Sarsa(λ), Q-Learning, ...);
- The correlation between the two representations does not need to be formalized.

Conclusions

Thesis results:

- Extended (Brafman et al. 2018) to the context of RL;
- Formalization of *RL for LTL_f / LDL_f goals*, devising of a solution and analysis of the advantages;
- Provided an implementation and given experimental evidence of the goodness of our approach.

Future works:

- Minimal low-level representation to tackle LTL_f / LDL_f goals;
- Optimize FLLOAT, enrich RLTG;
- Try the approach in several real world applications;
- Design of ad-hoc algorithms (e.g. "automata-aware" exploration);
- Extend the approach to the framework of Multi-Agent Systems.

Publication: <https://arxiv.org/abs/1807.06333>. Status: submitted