Titolo

Marco Favorito

May 22, 2018

Contents

Chapter 1

Introduction

Chapter 2

Preliminaries

In this chapter we describe the background knowledge required for this work. We introduce Markov Decision Process (MDP) and Non-Markovian Reward Decision Process (NMRDP), common formalisms in the context of Reinforcement Learning. We describe Linear Temporal Logic over finite traces (LTL_f) and Linear Dynamic Logic over finite traces (LDL_f), that we use for define temporal goal in a RL setting. Then, we describe an important result about RL for NMRDP with $\text{LTL}_f/\text{LDL}_f$rewards, that is the basis for this work.

2.1 Reinforcement Learning

2.1.1 MDP

2.1.2 NMRDP

2.2 LTL_f and LDL_f

2.2.1 Linear Temporal Logic for finite traces: LTL_f

2.2.2 Linear Dynamic Logic for finite traces: LDL_f

2.2.3 LTL_f and LDL_f translation to automata

2.3 RL for NMRDP with $\text{LTL}_f/\text{LDL}_f$rewards

Chapter 3

RL for $\text{LTL}_f/\text{LDL}_f$ Goals

Chapter 4

Automata-based Reward shaping

Chapter 5

Experiments

Chapter 6

Conclusions

Appendix A

FLLOAT

Appendix B

RLTG

Bibliography

[1] Ronen I. Brafman, Giuseppe De Giacomo, and Fabio Patrizi. Specifying non-markovian rewards in mdps using ldl on finite traces (preliminary version). *CoRR*, abs/1706.08100, 2017.